

# DeepLearning\_Fall2023\_HW6

November 11, 2023

## 1 Deep Learning Homework 6 (Spring 2023)

This code is provided for Deep Learning class (601.482/682) Homework 6. For ease of implementation, we recommend working entire in Google Colaboratory.

@Copyright Cong Gao, the Johns Hopkins University, cgao11@jhu.edu. Modifications made by Hongtao Wu, Suzanna Sia, Hao Ding, Keith Harrigian, and Yiqing Shen.

### 1.0.1 Imports

```
[2]: ## Mount Google Drive Data (If using Google Colaboratory)
try:
    from google.colab import drive
    drive.mount('/content/gdrive')
except:
    print("Mounting Failed.")
```

Mounted at /content/gdrive

```
[ ]: ## Standard Library
import os
import json

## External Libraries
import numpy as np
import torch
import torch.nn as nn
from torchvision import transforms
from torch.autograd import Variable
import torch.nn.functional as functional
from torch.utils.data import Dataset, DataLoader
from skimage import io
import matplotlib.pyplot as plt
```

## 2 Problem 1: Unsupervised Pre-training

### 2.0.1 Training Hyperparameters

These are recommended hyperparameters - please feel free to use what works for you. Batch size can be changed if it does not match your memory, please state your batch step\_size in your report.

Dataset is available at: [https://livejohnshopkins-my.sharepoint.com/:u:/g/personal/yshen92\\_jh\\_edu/EcTxWAXspEExyMg](https://livejohnshopkins-my.sharepoint.com/:u:/g/personal/yshen92_jh_edu/EcTxWAXspEExyMg)

```
[ ]: ## Batch Size
train_batch_size = 10
validation_batch_size = 10

## Learning Rate
learning_rate = 0.001

# Epochs (Consider setting high and implementing early stopping)
num_epochs = 100
```

### 2.0.2 Data Paths

```
[ ]: # General Data Directory ##TODO: Please fill in the appropriate directory
data_dir = "/content/gdrive/MyDrive/Colab Notebooks/MLDL_HW6/
↳DeepLearning_Fall2023_HW6_data/HW6_data"

## Segmentation + Colorization Paths
segmentation_data_dir = f"{data_dir}/segmentation/"
colorization_data_dir = f"{data_dir}/colorization/"

# Mask JSON
mask_json = f"{data_dir}/mapping.json"
```

### 2.0.3 Data Loaders

We have provided you with some preprocessing code for the images but you should feel free to modify the class however you please to support your training schema. In the very least, you will have to modify the dataloader to support loading of the colorization dataset.

```
[ ]: ## Image Transforms
img_transform = transforms.Compose([
    transforms.ToTensor(),
])

## Image Dataloader
class ImageDataset(Dataset):

    """
    ImageDataset
```

```

"""

def __init__(self,
              input_dir,
              op,
              mask_json_path,
              transforms=None):
    """
    ##TODO: Add support for colorization dataset

    Args:
        input_dir (str): Path to either colorization or segmentation
    ↪directory
        op (str): One of "train", "val", or "test" signifying the desired
    ↪split
        mask_json_path (str): Path to mapping.json file
        transforms (list or None): Image transformations to apply upon
    ↪loading.
    """
    self.transform = transforms
    self.op = op
    with open(mask_json_path, 'r') as f:
        self.mask = json.load(f)
    self.mask_num = len(self.mask) # There are 6 categories: grey, dark
    ↪grey, and black
    self.mask_value = [value for value in self.mask.values()]
    self.mask_value.sort()
    try:
        if self.op == 'train':
            self.data_dir = os.path.join(input_dir, 'train')
        elif self.op == 'val':
            self.data_dir = os.path.join(input_dir, 'validation')
        elif self.op == 'test':
            self.data_dir = os.path.join(input_dir, 'test')
    except ValueError:
        print('op should be either train, val or test!')

def __len__(self):
    """
    """
    return len(next(os.walk(self.data_dir))[1])

def __getitem__(self,
                idx):
    """

```

```

"""

## Load Image and Parse Properties
img_name = str(idx) + '_input.jpg'
mask_name = str(idx) + '_mask.png'
img = io.imread(os.path.join(self.data_dir, str(idx), img_name))
mask = io.imread(os.path.join(self.data_dir, str(idx), mask_name))
if len(mask.shape) == 2:
    h, w = mask.shape
elif len(mask.shape) == 3:
    h, w, c = mask.shape
## Convert grey-scale label to one-hot encoding
new_mask = np.zeros((h, w, self.mask_num))
for idx in range(self.mask_num):
    #if the mask has 3 dimension use this code
    new_mask[:, :, idx] = mask[:, :, 0] == self.mask_value[idx]
    #if the mask has 1 dimension use the code below
    #new_mask[:, :, idx] = mask == self.mask_value[idx]
## Transform image and mask
if self.transform:
    img, mask = self.img_transform(img, new_mask)
# ## Use dictionary to output
# sample = {'img': img, 'mask': mask}
# return sample
return img, mask

def img_transform(self,
                  img,
                  mask):
"""

"""

## Apply Transformations to Image and Mask
img = self.transform(img)
mask = self.transform(mask)
return img, mask

```

## 2.1 Model Architecture

Finish building the U-net architecture below.

```

[ ]: ## Functions for adding the convolution layer
def add_conv_stage(dim_in,
                  dim_out,
                  kernel_size=3,
                  stride=1,
                  padding=1,
                  bias=True,

```

```

        useBN=True):

    """

    """
    # Use batch normalization
    if useBN:
        return nn.Sequential(
            nn.Conv2d(dim_in, dim_out, kernel_size=kernel_size, stride=stride,
↳padding=padding, bias=bias),
            nn.BatchNorm2d(dim_out),
            nn.LeakyReLU(0.1),
            nn.Conv2d(dim_out, dim_out, kernel_size=kernel_size, stride=stride,
↳padding=padding, bias=bias),
            nn.BatchNorm2d(dim_out),
            nn.LeakyReLU(0.1)
        )
    # No batch normalization
    else:
        return nn.Sequential(
            nn.Conv2d(dim_in, dim_out, kernel_size=kernel_size, stride=stride,
↳padding=padding, bias=bias),
            nn.ReLU(),
            nn.Conv2d(dim_out, dim_out, kernel_size=kernel_size, stride=stride,
↳padding=padding, bias=bias),
            nn.ReLU()
        )

## Upsampling
def upsample(ch_coarse,
            ch_fine):
    """

    """
    return nn.Sequential(
        nn.ConvTranspose2d(ch_coarse, ch_fine, 4, 2, 1, bias=False),
        nn.ReLU())

# U-Net
class UNET(nn.Module):

    """

    """
    def __init__(self, n_classes, useBN=True):
        """
        Args:

```

```

        n_classes (int): Number of classes
        useBN (bool): Turn Batch Norm on or off. (Hint: Using BatchNorm_
↳ might help you achieve better performance.)
    """
    super(UNET, self).__init__()
    # Downgrade stages
    self.conv1 = add_conv_stage(3, 32, useBN=useBN)
    self.conv2 = add_conv_stage(32, 64, useBN=useBN)
    self.conv3 = add_conv_stage(64, 128, useBN=useBN)
    self.conv4 = add_conv_stage(128, 256, useBN=useBN)
    # Upgrade stages
    self.conv3m = add_conv_stage(256, 128, useBN=useBN)
    self.conv2m = add_conv_stage(128, 64, useBN=useBN)
    self.conv1m = add_conv_stage(64, 32, useBN=useBN)
    # Maxpool
    self.max_pool = nn.MaxPool2d(2)
    # Upsample layers
    self.upsample43 = upsample(256, 128)
    self.upsample32 = upsample(128, 64)
    self.upsample21 = upsample(64, 32)
    # weight initialization
    # You can have your own weight initialization. This is just an example.
    for m in self.modules():
        if isinstance(m, nn.Conv2d) or isinstance(m, nn.ConvTranspose2d):
            if m.bias is not None:
                m.bias.data.zero_()
    #TODO: Design your last layer & activations
    self.out = nn.Conv2d(32, n_classes, kernel_size = 1, stride = 1, padding_
↳ = 0, bias = True)
    self.softmax = nn.Softmax()

    def forward(self, x):
        """
        Forward pass
        """
        conv1_out = self.conv1(x)
        conv2_out = self.conv2(self.max_pool(conv1_out))
        conv3_out = self.conv3(self.max_pool(conv2_out))
        conv4_out = self.conv4(self.max_pool(conv3_out))

        conv4m_out_ = torch.cat((self.upsample43(conv4_out), conv3_out), 1)
        conv3m_out_ = self.conv3m(conv4m_out_)

        conv3m_out_ = torch.cat((self.upsample32(conv3m_out_), conv2_out), 1)
        conv2m_out_ = self.conv2m(conv3m_out_)

        conv2m_out_ = torch.cat((self.upsample21(conv2m_out_), conv1_out), 1)

```

```

conv1m_out = self.conv1m(conv2m_out_)

#TODO: Design your last layer & activations
out = self.softmax(self.out(conv1m_out))

return out

```

### 2.1.1 DICE Score and DICE Loss

Finish implementing the DICE score function below and then write a Dice Loss function that you can use to update your model weights.

```

[ ]: ##TODO: Finish implementing the multi-class DICE score function
def dice_score_image(prediction, target, n_classes):
    '''
        computer the mean dice score for a single image

        Reminders: A false positive is a result that indicates a given condition
        ↪exists, when it does not
                   A false negative is a test result that indicates that a
        ↪condition does not hold, while in fact it does

        Args:
            prediction (tensor): predicted labels of the image
            target (tensor): ground truth of the image
            n_classes (int): number of classes

        Returns:
            m_dice (float): Mean dice score over classes
    '''
    ## Should test image one by one
    assert prediction.shape[0] == 1 #This line can not be deleted
    ## TODO: Compute Dice Score for Each Class. Compute Mean Dice Score over
    ↪Classes.
    dice_classes = np.zeros(n_classes)
    for cl in range(n_classes):
        pred = (prediction[0] == cl).float()
        label = target[0, cl, :, :]

        TP = (pred * label).sum()
        FP = (pred * (1 - label)).sum()
        FN = ((1 - pred) * label).sum()
        #When there is no ground truth of the class in this image
        #Give 1 dice score if False Positive pixel number is 0,
        #give 0 dice score if False Positive pixel number is not 0 (> 0).
        if TP == 0 and FN == 0:
            if FP == 0:
                dice_classes[cl] = 1

```

```

        elif FP > 0:
            dice_classes[c1] = 0
        else:
            dice_classes[c1] = (2 * TP) / (2*TP + FP + FN)
    return dice_classes.mean()

def dice_score_dataset(model, dataloader, num_classes, use_gpu=False):
    """
    Compute the mean dice score on a set of data.

    Note that multiclass dice score can be defined as the mean over classes of
    ↪ binary
    ↪ dice score. Dice score is computed per image. Mean dice score over the
    ↪ dataset is the dice
    ↪ score averaged across all images.

    Reminders: A false positive is a result that indicates a given condition
    ↪ exists, when it does not
    ↪ A false negative is a test result that indicates that a
    ↪ condition does not hold, while in fact it does

    Args:
        model (UNET class): Your trained model
        dataloader (DataLoader): Dataset for evaluation
        num_classes (int): Number of classes

    Returns:
        m_dice (float): Mean dice score over the input dataset
    """
    ## Number of Batches and Cache over Dataset
    n_batches = len(dataloader)
    scores = np.zeros(n_batches)
    ## Evaluate
    model.eval()
    idx = 0
    for data in dataloader:
        ## Format Data
        img, target = data
        if use_gpu:
            img = img.cuda()
            target = target.cuda()
        ## Make Predictions
        out = model(img)
        n_classes = out.shape[1]
        prediction = torch.argmax(out, dim = 1)

```



```

        scores[idx] = dice_score_image(prediction, target, n_classes)
        idx += 1
    ## Average Dice Score Over Images
    m_dice = scores.mean()
    return m_dice

## TODO: Implement DICE loss,
# It should conform to to how we computer the dice score.
class DICELoss(nn.Module):
    def __init__(self, num_classes):
        super(DICELoss, self).__init__()
        self.num_classes = num_classes
    def forward(self, input, target):
        intersection = 0
        pred_sum = 0
        label_sum = 0
        for i in range(input.shape[0]):
            for cl in range(self.num_classes):
                pred = input[i, cl, :, :]
                label = target[i, cl, :, :]
                intersection += (pred * label).sum()
                pred_sum += (pred * pred).sum()
                label_sum += (label * label).sum()
        total_loss = 1 - (2 * intersection)/(pred_sum + label_sum)
        return total_loss

```

## 2.2 Training Procedure (Segmentation)

```

[ ]: ## Initialize your unet
n_classes = 6
model = UNET(n_classes)

## Initialize Dataloaders
train_dataset=ImageDataset(input_dir=segmentation_data_dir, op="train",
    ↪mask_json_path=mask_json, transforms=img_transform)
validation_dataset=ImageDataset(input_dir=segmentation_data_dir, op="val",
    ↪mask_json_path=mask_json, transforms=img_transform)
test_dataset=ImageDataset(input_dir=segmentation_data_dir, op="test",
    ↪mask_json_path=mask_json, transforms=img_transform)
train_dataloader = DataLoader(train_dataset, batch_size=train_batch_size,
    ↪shuffle=True)
validation_dataloader = DataLoader(validation_dataset,
    ↪batch_size=validation_batch_size, shuffle=False)
test_dataloader = DataLoader(test_dataset, batch_size=1, shuffle=False)
## Initialize Optimizer and Learning Rate Scheduler

```

```

optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=10, gamma=0.1)
criterion = DICELoss(n_classes)
model = model.cuda()

# Define lists to store training and validation losses
train_losses = []
val_losses = []
test_score = []

print("Start Training...")
for epoch in range(num_epochs):
    ##### Training #####
    print("\nEPOCH " + str(epoch+1) + " of " + str(num_epochs) + "\n")
    # TODO: Design your own training section
    model.train()
    train_loss = 0.0

    # Iterate through the training data
    for inputs, labels in train_dataloader:
        inputs = inputs.cuda()
        labels = labels.cuda()
        output = model.forward(inputs)
        loss = criterion.forward(output, labels)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        train_loss += loss.item()

    avg_train_loss = train_loss / len(train_dataloader)
    train_losses.append(avg_train_loss)

    ##### Validation #####
    # TODO: Design your own validation section
    model.eval()
    val_loss = 0.0
    with torch.no_grad():
        for inputs, labels in validation_dataloader:
            inputs = inputs.cuda()
            labels = labels.cuda()
            output = model(inputs)
            loss = criterion.forward(output, labels)
            val_loss += loss.item()

    avg_val_loss = val_loss / len(validation_dataloader)

```

```

    val_losses.append(avg_val_loss)
    test_score_i = dice_score_dataset(model, test_dataloader, n_classes,
↪use_gpu=gpu_bool)
    test_score.append(test_score_i)
    scheduler.step()

    print(f'Epoch [{epoch+1}/{num_epochs}], Train Loss: {avg_train_loss:.4f},
↪Val Loss: {avg_val_loss:.4f}')
    print(f'Test Score: {test_score_i:.4f}')

```

Start Training...

EPOCH 1 of 100

```

/usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py:1518:
UserWarning: Implicit dimension choice for softmax has been deprecated. Change
the call to include dim=X as an argument.
    return self._call_impl(*args, **kwargs)

```

Epoch [1/100], Train Loss: 0.3988, Val Loss: 0.2168  
Test Score: 0.5909

EPOCH 2 of 100

Epoch [2/100], Train Loss: 0.2262, Val Loss: 0.1446  
Test Score: 0.6099

EPOCH 3 of 100

Epoch [3/100], Train Loss: 0.1499, Val Loss: 0.1124  
Test Score: 0.5028

EPOCH 4 of 100

Epoch [4/100], Train Loss: 0.1098, Val Loss: 0.1014  
Test Score: 0.5022

EPOCH 5 of 100

Epoch [5/100], Train Loss: 0.0947, Val Loss: 0.1183  
Test Score: 0.5234

EPOCH 6 of 100

Epoch [6/100], Train Loss: 0.0824, Val Loss: 0.0766  
Test Score: 0.4657

EPOCH 7 of 100

Epoch [7/100], Train Loss: 0.0756, Val Loss: 0.1432  
Test Score: 0.3881

EPOCH 8 of 100

Epoch [8/100], Train Loss: 0.0725, Val Loss: 0.0766  
Test Score: 0.5347

EPOCH 9 of 100

Epoch [9/100], Train Loss: 0.0698, Val Loss: 0.0696  
Test Score: 0.5196

EPOCH 10 of 100

Epoch [10/100], Train Loss: 0.0650, Val Loss: 0.0701  
Test Score: 0.5394

EPOCH 11 of 100

Epoch [11/100], Train Loss: 0.0599, Val Loss: 0.0613  
Test Score: 0.5804

EPOCH 12 of 100

Epoch [12/100], Train Loss: 0.0593, Val Loss: 0.0591  
Test Score: 0.5667

EPOCH 13 of 100

Epoch [13/100], Train Loss: 0.0575, Val Loss: 0.0580  
Test Score: 0.6237

EPOCH 14 of 100

Epoch [14/100], Train Loss: 0.0572, Val Loss: 0.0571  
Test Score: 0.6174

EPOCH 15 of 100

Epoch [15/100], Train Loss: 0.0549, Val Loss: 0.0566  
Test Score: 0.6213

EPOCH 16 of 100

Epoch [16/100], Train Loss: 0.0574, Val Loss: 0.0581

Test Score: 0.6246

EPOCH 17 of 100

Epoch [17/100], Train Loss: 0.0554, Val Loss: 0.0550  
Test Score: 0.6139

EPOCH 18 of 100

Epoch [18/100], Train Loss: 0.0543, Val Loss: 0.0549  
Test Score: 0.6104

EPOCH 19 of 100

Epoch [19/100], Train Loss: 0.0534, Val Loss: 0.0578  
Test Score: 0.6248

EPOCH 20 of 100

Epoch [20/100], Train Loss: 0.0535, Val Loss: 0.0541  
Test Score: 0.6440

EPOCH 21 of 100

Epoch [21/100], Train Loss: 0.0500, Val Loss: 0.0524  
Test Score: 0.6092

EPOCH 22 of 100

Epoch [22/100], Train Loss: 0.0506, Val Loss: 0.0522  
Test Score: 0.6168

EPOCH 23 of 100

Epoch [23/100], Train Loss: 0.0506, Val Loss: 0.0524  
Test Score: 0.6118

EPOCH 24 of 100

Epoch [24/100], Train Loss: 0.0512, Val Loss: 0.0524  
Test Score: 0.6090

EPOCH 25 of 100

Epoch [25/100], Train Loss: 0.0519, Val Loss: 0.0522  
Test Score: 0.6073

EPOCH 26 of 100

Epoch [26/100], Train Loss: 0.0517, Val Loss: 0.0522  
Test Score: 0.6151

EPOCH 27 of 100

Epoch [27/100], Train Loss: 0.0492, Val Loss: 0.0522  
Test Score: 0.6126

EPOCH 28 of 100

Epoch [28/100], Train Loss: 0.0525, Val Loss: 0.0521  
Test Score: 0.6336

EPOCH 29 of 100

Epoch [29/100], Train Loss: 0.0493, Val Loss: 0.0517  
Test Score: 0.6029

EPOCH 30 of 100

Epoch [30/100], Train Loss: 0.0505, Val Loss: 0.0519  
Test Score: 0.6064

EPOCH 31 of 100

Epoch [31/100], Train Loss: 0.0493, Val Loss: 0.0517  
Test Score: 0.6055

EPOCH 32 of 100

Epoch [32/100], Train Loss: 0.0491, Val Loss: 0.0521  
Test Score: 0.6006

EPOCH 33 of 100

Epoch [33/100], Train Loss: 0.0510, Val Loss: 0.0520  
Test Score: 0.6035

EPOCH 34 of 100

Epoch [34/100], Train Loss: 0.0491, Val Loss: 0.0520  
Test Score: 0.6163

EPOCH 35 of 100

Epoch [35/100], Train Loss: 0.0495, Val Loss: 0.0519  
Test Score: 0.6002

EPOCH 36 of 100

Epoch [36/100], Train Loss: 0.0520, Val Loss: 0.0516  
Test Score: 0.6207

EPOCH 37 of 100

Epoch [37/100], Train Loss: 0.0504, Val Loss: 0.0517  
Test Score: 0.6203

EPOCH 38 of 100

Epoch [38/100], Train Loss: 0.0512, Val Loss: 0.0517  
Test Score: 0.6033

EPOCH 39 of 100

Epoch [39/100], Train Loss: 0.0517, Val Loss: 0.0517  
Test Score: 0.6101

EPOCH 40 of 100

Epoch [40/100], Train Loss: 0.0504, Val Loss: 0.0518  
Test Score: 0.5999

EPOCH 41 of 100

Epoch [41/100], Train Loss: 0.0510, Val Loss: 0.0519  
Test Score: 0.6155

EPOCH 42 of 100

Epoch [42/100], Train Loss: 0.0497, Val Loss: 0.0515  
Test Score: 0.6091

EPOCH 43 of 100

Epoch [43/100], Train Loss: 0.0501, Val Loss: 0.0521  
Test Score: 0.5959

EPOCH 44 of 100

Epoch [44/100], Train Loss: 0.0508, Val Loss: 0.0521  
Test Score: 0.6036

EPOCH 45 of 100

Epoch [45/100], Train Loss: 0.0494, Val Loss: 0.0517  
Test Score: 0.5968

EPOCH 46 of 100

Epoch [46/100], Train Loss: 0.0503, Val Loss: 0.0518  
Test Score: 0.5967

EPOCH 47 of 100

Epoch [47/100], Train Loss: 0.0504, Val Loss: 0.0517  
Test Score: 0.6138

EPOCH 48 of 100

Epoch [48/100], Train Loss: 0.0503, Val Loss: 0.0518  
Test Score: 0.6032

EPOCH 49 of 100

Epoch [49/100], Train Loss: 0.0492, Val Loss: 0.0517  
Test Score: 0.6034

EPOCH 50 of 100

Epoch [50/100], Train Loss: 0.0492, Val Loss: 0.0519  
Test Score: 0.6073

EPOCH 51 of 100

Epoch [51/100], Train Loss: 0.0507, Val Loss: 0.0518  
Test Score: 0.5897

EPOCH 52 of 100

Epoch [52/100], Train Loss: 0.0490, Val Loss: 0.0518  
Test Score: 0.6003

EPOCH 53 of 100

Epoch [53/100], Train Loss: 0.0504, Val Loss: 0.0516  
Test Score: 0.6128

EPOCH 54 of 100

Epoch [54/100], Train Loss: 0.0497, Val Loss: 0.0517  
Test Score: 0.6129



EPOCH 55 of 100

Epoch [55/100], Train Loss: 0.0509, Val Loss: 0.0517  
Test Score: 0.6064

EPOCH 56 of 100

Epoch [56/100], Train Loss: 0.0503, Val Loss: 0.0516  
Test Score: 0.6030

EPOCH 57 of 100

Epoch [57/100], Train Loss: 0.0506, Val Loss: 0.0517  
Test Score: 0.6030

EPOCH 58 of 100

Epoch [58/100], Train Loss: 0.0500, Val Loss: 0.0517  
Test Score: 0.6060

EPOCH 59 of 100

Epoch [59/100], Train Loss: 0.0517, Val Loss: 0.0516  
Test Score: 0.6092

EPOCH 60 of 100

Epoch [60/100], Train Loss: 0.0502, Val Loss: 0.0516  
Test Score: 0.6093

EPOCH 61 of 100

Epoch [61/100], Train Loss: 0.0497, Val Loss: 0.0515  
Test Score: 0.6122

EPOCH 62 of 100

Epoch [62/100], Train Loss: 0.0500, Val Loss: 0.0518  
Test Score: 0.6133

EPOCH 63 of 100

Epoch [63/100], Train Loss: 0.0498, Val Loss: 0.0519  
Test Score: 0.5969

EPOCH 64 of 100

Epoch [64/100], Train Loss: 0.0504, Val Loss: 0.0518

Test Score: 0.6036

EPOCH 65 of 100

Epoch [65/100], Train Loss: 0.0500, Val Loss: 0.0519  
Test Score: 0.6131

EPOCH 66 of 100

Epoch [66/100], Train Loss: 0.0489, Val Loss: 0.0518  
Test Score: 0.6069

EPOCH 67 of 100

Epoch [67/100], Train Loss: 0.0491, Val Loss: 0.0518  
Test Score: 0.6035

EPOCH 68 of 100

Epoch [68/100], Train Loss: 0.0504, Val Loss: 0.0517  
Test Score: 0.6063

EPOCH 69 of 100

Epoch [69/100], Train Loss: 0.0508, Val Loss: 0.0519  
Test Score: 0.5964

EPOCH 70 of 100

Epoch [70/100], Train Loss: 0.0496, Val Loss: 0.0516  
Test Score: 0.6061

EPOCH 71 of 100

Epoch [71/100], Train Loss: 0.0498, Val Loss: 0.0517  
Test Score: 0.5927

EPOCH 72 of 100

Epoch [72/100], Train Loss: 0.0488, Val Loss: 0.0518  
Test Score: 0.6033

EPOCH 73 of 100

Epoch [73/100], Train Loss: 0.0505, Val Loss: 0.0519  
Test Score: 0.6311

EPOCH 74 of 100

Epoch [74/100], Train Loss: 0.0498, Val Loss: 0.0519  
Test Score: 0.6003

EPOCH 75 of 100

Epoch [75/100], Train Loss: 0.0500, Val Loss: 0.0519  
Test Score: 0.5967

EPOCH 76 of 100

Epoch [76/100], Train Loss: 0.0510, Val Loss: 0.0518  
Test Score: 0.6031

EPOCH 77 of 100

Epoch [77/100], Train Loss: 0.0521, Val Loss: 0.0517  
Test Score: 0.6031

EPOCH 78 of 100

Epoch [78/100], Train Loss: 0.0506, Val Loss: 0.0520  
Test Score: 0.6076

EPOCH 79 of 100

Epoch [79/100], Train Loss: 0.0494, Val Loss: 0.0519  
Test Score: 0.6005

EPOCH 80 of 100

Epoch [80/100], Train Loss: 0.0519, Val Loss: 0.0516  
Test Score: 0.6126

EPOCH 81 of 100

Epoch [81/100], Train Loss: 0.0505, Val Loss: 0.0519  
Test Score: 0.6056

EPOCH 82 of 100

Epoch [82/100], Train Loss: 0.0493, Val Loss: 0.0518  
Test Score: 0.6069

EPOCH 83 of 100

Epoch [83/100], Train Loss: 0.0512, Val Loss: 0.0517  
Test Score: 0.6065

EPOCH 84 of 100

Epoch [84/100], Train Loss: 0.0507, Val Loss: 0.0517  
Test Score: 0.6032

EPOCH 85 of 100

Epoch [85/100], Train Loss: 0.0507, Val Loss: 0.0517  
Test Score: 0.6037

EPOCH 86 of 100

Epoch [86/100], Train Loss: 0.0504, Val Loss: 0.0519  
Test Score: 0.6106

EPOCH 87 of 100

Epoch [87/100], Train Loss: 0.0503, Val Loss: 0.0519  
Test Score: 0.6068

EPOCH 88 of 100

Epoch [88/100], Train Loss: 0.0499, Val Loss: 0.0523  
Test Score: 0.6107

EPOCH 89 of 100

Epoch [89/100], Train Loss: 0.0494, Val Loss: 0.0516  
Test Score: 0.6187

EPOCH 90 of 100

Epoch [90/100], Train Loss: 0.0499, Val Loss: 0.0517  
Test Score: 0.6035

EPOCH 91 of 100

Epoch [91/100], Train Loss: 0.0503, Val Loss: 0.0516  
Test Score: 0.6124

EPOCH 92 of 100

Epoch [92/100], Train Loss: 0.0500, Val Loss: 0.0516  
Test Score: 0.5962

EPOCH 93 of 100

Epoch [93/100], Train Loss: 0.0496, Val Loss: 0.0518  
Test Score: 0.6007

EPOCH 94 of 100

Epoch [94/100], Train Loss: 0.0492, Val Loss: 0.0516  
Test Score: 0.6095

EPOCH 95 of 100

Epoch [95/100], Train Loss: 0.0501, Val Loss: 0.0517  
Test Score: 0.6031

EPOCH 96 of 100

Epoch [96/100], Train Loss: 0.0507, Val Loss: 0.0518  
Test Score: 0.6031

EPOCH 97 of 100

Epoch [97/100], Train Loss: 0.0502, Val Loss: 0.0518  
Test Score: 0.5996

EPOCH 98 of 100

Epoch [98/100], Train Loss: 0.0511, Val Loss: 0.0517  
Test Score: 0.6126

EPOCH 99 of 100

Epoch [99/100], Train Loss: 0.0507, Val Loss: 0.0516  
Test Score: 0.6065

EPOCH 100 of 100

Epoch [100/100], Train Loss: 0.0498, Val Loss: 0.0516  
Test Score: 0.6192

```
[ ]: max_score = max(test_score)
      max_score_idx = test_score.index(max_score)
      print(f'At Epoch {max_score_idx}, Test DICE score reaches its max as:␣
            ↳{max_score:.4f}')
```

At Epoch 19, Test DICE score reaches its max as: 0.6440

```
[ ]: # Graph
      fig, axs = plt.subplots(1, 2,figsize=(20, 7))
```

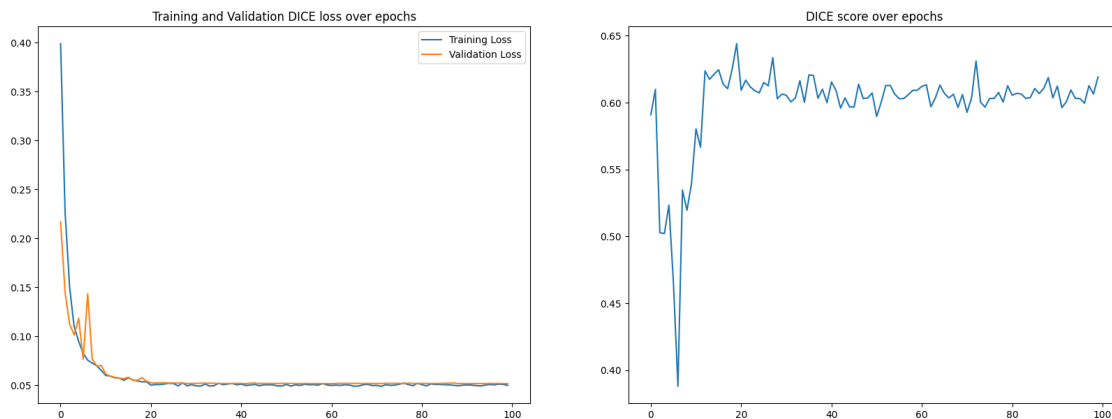
```

axs[0].plot(range(num_epochs), train_losses)
axs[0].plot(range(num_epochs), val_losses)
axs[0].legend(['Training Loss', 'Validation Loss'])
axs[0].set_title('Training and Validation DICE loss over epochs')

axs[1].plot(range(num_epochs), test_score)
axs[1].set_title('DICE score over epochs')

plt.show()

```



## 2.3 Data Augmentation

```

[ ]: ## Augmented Image Transforms
augmented_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.RandomHorizontalFlip(),
    transforms.RandomVerticalFlip(),
])

[ ]: ## Initialize your unet
n_classes = 6
model = UNET(n_classes)

## Initialize Dataloaders
train_dataset=ImageDataset(input_dir=segmentation_data_dir, op="train",
    ↪mask_json_path=mask_json, transforms=augmented_transform)
validation_dataset=ImageDataset(input_dir=segmentation_data_dir, op="val",
    ↪mask_json_path=mask_json, transforms=img_transform)
test_dataset=ImageDataset(input_dir=segmentation_data_dir, op="test",
    ↪mask_json_path=mask_json, transforms=img_transform)

```

```

train_dataloader = DataLoader(train_dataset, batch_size=train_batch_size,
    ↪shuffle=True)
validation_dataloader = DataLoader(validation_dataset,
    ↪batch_size=validation_batch_size, shuffle=False)
test_dataloader = DataLoader(test_dataset, batch_size=1, shuffle=False)
## Initialize Optimizer and Learning Rate Scheduler
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=10, gamma=0.1)
criterion = DICELoss(n_classes)
model = model.cuda()

# Define lists to store training and validation losses
train_losses = []
val_losses = []
test_score = []

print("Start Training...")
for epoch in range(num_epochs):
    ##### Training #####
    print("\nEPOCH " + str(epoch+1) + " of " + str(num_epochs) + "\n")
    # TODO: Design your own training section
    model.train()
    train_loss = 0.0

    # Iterate through the training data
    for inputs, labels in train_dataloader:
        # Forward pass
        inputs = inputs.cuda()
        labels = labels.cuda()
        output = model.forward(inputs)

        # Compute loss
        loss = criterion.forward(output, labels)

        # Backpropagation
        optimizer.zero_grad() # Zero the gradients
        loss.backward()
        optimizer.step()

    # Update running training loss
    train_loss += loss.item()

    # Compute average training loss for the epoch
    avg_train_loss = train_loss / len(train_dataloader)
    train_losses.append(avg_train_loss)

    ##### Validation #####

```

```

# TODO: Design your own validation section
model.eval()
val_loss = 0.0
with torch.no_grad():
    for inputs, labels in validation_dataloader:
        if gpu_bool:
            inputs = inputs.cuda()
            labels = labels.cuda()
        output = model(inputs)
        loss = criterion.forward(output, labels)
        val_loss += loss.item()

avg_val_loss = val_loss / len(validation_dataloader)
val_losses.append(avg_val_loss)

##### Test #####
test_score_i = dice_score_dataset(model, test_dataloader, n_classes,
↪use_gpu=gpu_bool)
test_score.append(test_score_i)
scheduler.step()

print(f'Epoch [{epoch+1}/{num_epochs}], Train Loss: {avg_train_loss:.4f},
↪Val Loss: {avg_val_loss:.4f}')
print(f'Test Score: {test_score_i:.4f}')

```

Start Training...

EPOCH 1 of 100

Epoch [1/100], Train Loss: 0.5144, Val Loss: 0.3952  
Test Score: 0.3992

EPOCH 2 of 100

Epoch [2/100], Train Loss: 0.3254, Val Loss: 0.3205  
Test Score: 0.2604

EPOCH 3 of 100

Epoch [3/100], Train Loss: 0.2249, Val Loss: 0.2182  
Test Score: 0.5009

EPOCH 4 of 100

Epoch [4/100], Train Loss: 0.1820, Val Loss: 0.1834  
Test Score: 0.5757



EPOCH 5 of 100

Epoch [5/100], Train Loss: 0.1600, Val Loss: 0.1823  
Test Score: 0.3985

EPOCH 6 of 100

Epoch [6/100], Train Loss: 0.1547, Val Loss: 0.1637  
Test Score: 0.6829

EPOCH 7 of 100

Epoch [7/100], Train Loss: 0.1480, Val Loss: 0.1847  
Test Score: 0.5854

EPOCH 8 of 100

Epoch [8/100], Train Loss: 0.1455, Val Loss: 0.1583  
Test Score: 0.6540

EPOCH 9 of 100

Epoch [9/100], Train Loss: 0.1440, Val Loss: 0.1566  
Test Score: 0.6669

EPOCH 10 of 100

Epoch [10/100], Train Loss: 0.1421, Val Loss: 0.1551  
Test Score: 0.6559

EPOCH 11 of 100

Epoch [11/100], Train Loss: 0.1406, Val Loss: 0.1518  
Test Score: 0.6799

EPOCH 12 of 100

Epoch [12/100], Train Loss: 0.1410, Val Loss: 0.1542  
Test Score: 0.6701

EPOCH 13 of 100

Epoch [13/100], Train Loss: 0.1406, Val Loss: 0.1553  
Test Score: 0.6701

EPOCH 14 of 100

Epoch [14/100], Train Loss: 0.1391, Val Loss: 0.1520

Test Score: 0.6600

EPOCH 15 of 100

Epoch [15/100], Train Loss: 0.1411, Val Loss: 0.1574  
Test Score: 0.6703

EPOCH 16 of 100

Epoch [16/100], Train Loss: 0.1394, Val Loss: 0.1548  
Test Score: 0.6677

EPOCH 17 of 100

Epoch [17/100], Train Loss: 0.1405, Val Loss: 0.1509  
Test Score: 0.6734

EPOCH 18 of 100

Epoch [18/100], Train Loss: 0.1376, Val Loss: 0.1555  
Test Score: 0.6511

EPOCH 19 of 100

Epoch [19/100], Train Loss: 0.1415, Val Loss: 0.1554  
Test Score: 0.6798

EPOCH 20 of 100

Epoch [20/100], Train Loss: 0.1419, Val Loss: 0.1535  
Test Score: 0.6235

EPOCH 21 of 100

Epoch [21/100], Train Loss: 0.1404, Val Loss: 0.1539  
Test Score: 0.6333

EPOCH 22 of 100

Epoch [22/100], Train Loss: 0.1391, Val Loss: 0.1536  
Test Score: 0.6467

EPOCH 23 of 100

Epoch [23/100], Train Loss: 0.1399, Val Loss: 0.1486  
Test Score: 0.6483

EPOCH 24 of 100

Epoch [24/100], Train Loss: 0.1394, Val Loss: 0.1604  
Test Score: 0.6679

EPOCH 25 of 100

Epoch [25/100], Train Loss: 0.1410, Val Loss: 0.1508  
Test Score: 0.6680

EPOCH 26 of 100

Epoch [26/100], Train Loss: 0.1404, Val Loss: 0.1492  
Test Score: 0.6580

EPOCH 27 of 100

Epoch [27/100], Train Loss: 0.1385, Val Loss: 0.1520  
Test Score: 0.6566

EPOCH 28 of 100

Epoch [28/100], Train Loss: 0.1395, Val Loss: 0.1549  
Test Score: 0.6643

EPOCH 29 of 100

Epoch [29/100], Train Loss: 0.1402, Val Loss: 0.1514  
Test Score: 0.6641

EPOCH 30 of 100

Epoch [30/100], Train Loss: 0.1392, Val Loss: 0.1567  
Test Score: 0.6673

EPOCH 31 of 100

Epoch [31/100], Train Loss: 0.1393, Val Loss: 0.1559  
Test Score: 0.6536

EPOCH 32 of 100

Epoch [32/100], Train Loss: 0.1406, Val Loss: 0.1524  
Test Score: 0.6635

EPOCH 33 of 100

Epoch [33/100], Train Loss: 0.1373, Val Loss: 0.1540  
Test Score: 0.6704

EPOCH 34 of 100

Epoch [34/100], Train Loss: 0.1403, Val Loss: 0.1515  
Test Score: 0.6732

EPOCH 35 of 100

Epoch [35/100], Train Loss: 0.1398, Val Loss: 0.1504  
Test Score: 0.6767

EPOCH 36 of 100

Epoch [36/100], Train Loss: 0.1394, Val Loss: 0.1508  
Test Score: 0.6602

EPOCH 37 of 100

Epoch [37/100], Train Loss: 0.1387, Val Loss: 0.1566  
Test Score: 0.6700

EPOCH 38 of 100

Epoch [38/100], Train Loss: 0.1378, Val Loss: 0.1513  
Test Score: 0.6673

EPOCH 39 of 100

Epoch [39/100], Train Loss: 0.1391, Val Loss: 0.1492  
Test Score: 0.6701

EPOCH 40 of 100

Epoch [40/100], Train Loss: 0.1400, Val Loss: 0.1471  
Test Score: 0.6665

EPOCH 41 of 100

Epoch [41/100], Train Loss: 0.1400, Val Loss: 0.1502  
Test Score: 0.6666

EPOCH 42 of 100

Epoch [42/100], Train Loss: 0.1390, Val Loss: 0.1584  
Test Score: 0.6598

EPOCH 43 of 100

Epoch [43/100], Train Loss: 0.1407, Val Loss: 0.1563  
Test Score: 0.6697

EPOCH 44 of 100

Epoch [44/100], Train Loss: 0.1399, Val Loss: 0.1576  
Test Score: 0.6644

EPOCH 45 of 100

Epoch [45/100], Train Loss: 0.1407, Val Loss: 0.1572  
Test Score: 0.6698

EPOCH 46 of 100

Epoch [46/100], Train Loss: 0.1395, Val Loss: 0.1528  
Test Score: 0.6636

EPOCH 47 of 100

Epoch [47/100], Train Loss: 0.1381, Val Loss: 0.1553  
Test Score: 0.6705

EPOCH 48 of 100

Epoch [48/100], Train Loss: 0.1390, Val Loss: 0.1507  
Test Score: 0.6675

EPOCH 49 of 100

Epoch [49/100], Train Loss: 0.1391, Val Loss: 0.1515  
Test Score: 0.6636

EPOCH 50 of 100

Epoch [50/100], Train Loss: 0.1391, Val Loss: 0.1516  
Test Score: 0.6669

EPOCH 51 of 100

Epoch [51/100], Train Loss: 0.1384, Val Loss: 0.1549  
Test Score: 0.6653

EPOCH 52 of 100

Epoch [52/100], Train Loss: 0.1398, Val Loss: 0.1525  
Test Score: 0.6634

EPOCH 53 of 100

Epoch [53/100], Train Loss: 0.1402, Val Loss: 0.1559  
Test Score: 0.6641

EPOCH 54 of 100

Epoch [54/100], Train Loss: 0.1392, Val Loss: 0.1510  
Test Score: 0.6606

EPOCH 55 of 100

Epoch [55/100], Train Loss: 0.1386, Val Loss: 0.1508  
Test Score: 0.6715

EPOCH 56 of 100

Epoch [56/100], Train Loss: 0.1395, Val Loss: 0.1475  
Test Score: 0.6668

EPOCH 57 of 100

Epoch [57/100], Train Loss: 0.1377, Val Loss: 0.1540  
Test Score: 0.6580

EPOCH 58 of 100

Epoch [58/100], Train Loss: 0.1399, Val Loss: 0.1584  
Test Score: 0.6643

EPOCH 59 of 100

Epoch [59/100], Train Loss: 0.1383, Val Loss: 0.1559  
Test Score: 0.6629

EPOCH 60 of 100

Epoch [60/100], Train Loss: 0.1372, Val Loss: 0.1537  
Test Score: 0.6566

EPOCH 61 of 100

Epoch [61/100], Train Loss: 0.1401, Val Loss: 0.1567  
Test Score: 0.6570

EPOCH 62 of 100

Epoch [62/100], Train Loss: 0.1400, Val Loss: 0.1546

Test Score: 0.6709

EPOCH 63 of 100

Epoch [63/100], Train Loss: 0.1398, Val Loss: 0.1536  
Test Score: 0.6732

EPOCH 64 of 100

Epoch [64/100], Train Loss: 0.1411, Val Loss: 0.1531  
Test Score: 0.6607

EPOCH 65 of 100

Epoch [65/100], Train Loss: 0.1378, Val Loss: 0.1501  
Test Score: 0.6703

EPOCH 66 of 100

Epoch [66/100], Train Loss: 0.1383, Val Loss: 0.1531  
Test Score: 0.6677

EPOCH 67 of 100

Epoch [67/100], Train Loss: 0.1385, Val Loss: 0.1567  
Test Score: 0.6700

EPOCH 68 of 100

Epoch [68/100], Train Loss: 0.1385, Val Loss: 0.1506  
Test Score: 0.6638

EPOCH 69 of 100

Epoch [69/100], Train Loss: 0.1403, Val Loss: 0.1521  
Test Score: 0.6678

EPOCH 70 of 100

Epoch [70/100], Train Loss: 0.1395, Val Loss: 0.1505  
Test Score: 0.6603

EPOCH 71 of 100

Epoch [71/100], Train Loss: 0.1380, Val Loss: 0.1543  
Test Score: 0.6596

EPOCH 72 of 100

Epoch [72/100], Train Loss: 0.1396, Val Loss: 0.1556  
Test Score: 0.6636

EPOCH 73 of 100

Epoch [73/100], Train Loss: 0.1384, Val Loss: 0.1587  
Test Score: 0.6675

EPOCH 74 of 100

Epoch [74/100], Train Loss: 0.1374, Val Loss: 0.1495  
Test Score: 0.6533

EPOCH 75 of 100

Epoch [75/100], Train Loss: 0.1387, Val Loss: 0.1508  
Test Score: 0.6669

EPOCH 76 of 100

Epoch [76/100], Train Loss: 0.1403, Val Loss: 0.1502  
Test Score: 0.6702

EPOCH 77 of 100

Epoch [77/100], Train Loss: 0.1405, Val Loss: 0.1523  
Test Score: 0.6467

EPOCH 78 of 100

Epoch [78/100], Train Loss: 0.1384, Val Loss: 0.1482  
Test Score: 0.6610

EPOCH 79 of 100

Epoch [79/100], Train Loss: 0.1393, Val Loss: 0.1573  
Test Score: 0.6650

EPOCH 80 of 100

Epoch [80/100], Train Loss: 0.1385, Val Loss: 0.1554  
Test Score: 0.6576

EPOCH 81 of 100

Epoch [81/100], Train Loss: 0.1386, Val Loss: 0.1518  
Test Score: 0.6704



EPOCH 82 of 100

Epoch [82/100], Train Loss: 0.1410, Val Loss: 0.1486  
Test Score: 0.6712

EPOCH 83 of 100

Epoch [83/100], Train Loss: 0.1409, Val Loss: 0.1570  
Test Score: 0.6535

EPOCH 84 of 100

Epoch [84/100], Train Loss: 0.1395, Val Loss: 0.1523  
Test Score: 0.6696

EPOCH 85 of 100

Epoch [85/100], Train Loss: 0.1416, Val Loss: 0.1516  
Test Score: 0.6572

EPOCH 86 of 100

Epoch [86/100], Train Loss: 0.1398, Val Loss: 0.1508  
Test Score: 0.6706

EPOCH 87 of 100

Epoch [87/100], Train Loss: 0.1406, Val Loss: 0.1528  
Test Score: 0.6712

EPOCH 88 of 100

Epoch [88/100], Train Loss: 0.1385, Val Loss: 0.1515  
Test Score: 0.6700

EPOCH 89 of 100

Epoch [89/100], Train Loss: 0.1389, Val Loss: 0.1535  
Test Score: 0.6681

EPOCH 90 of 100

Epoch [90/100], Train Loss: 0.1404, Val Loss: 0.1554  
Test Score: 0.6569

EPOCH 91 of 100

Epoch [91/100], Train Loss: 0.1394, Val Loss: 0.1560  
Test Score: 0.6467

EPOCH 92 of 100

Epoch [92/100], Train Loss: 0.1380, Val Loss: 0.1498  
Test Score: 0.6702

EPOCH 93 of 100

Epoch [93/100], Train Loss: 0.1394, Val Loss: 0.1506  
Test Score: 0.6635

EPOCH 94 of 100

Epoch [94/100], Train Loss: 0.1377, Val Loss: 0.1497  
Test Score: 0.6675

EPOCH 95 of 100

Epoch [95/100], Train Loss: 0.1398, Val Loss: 0.1567  
Test Score: 0.6702

EPOCH 96 of 100

Epoch [96/100], Train Loss: 0.1411, Val Loss: 0.1540  
Test Score: 0.6668

EPOCH 97 of 100

Epoch [97/100], Train Loss: 0.1394, Val Loss: 0.1583  
Test Score: 0.6673

EPOCH 98 of 100

Epoch [98/100], Train Loss: 0.1386, Val Loss: 0.1513  
Test Score: 0.6604

EPOCH 99 of 100

Epoch [99/100], Train Loss: 0.1389, Val Loss: 0.1479  
Test Score: 0.6509

EPOCH 100 of 100

Epoch [100/100], Train Loss: 0.1406, Val Loss: 0.1552  
Test Score: 0.6599

```
[ ]: max_score = max(test_score)
max_score_idx = test_score.index(max_score)
print(f'At Epoch {max_score_idx}, Test DICE score reaches its max as:␣
↪{max_score:.4f}')
```


At Epoch 75, Test DICE score reaches its max as: 0.6831

```
[ ]: # Graph
fig, axs = plt.subplots(1, 2, figsize=(20, 7))

axs[0].plot(range(num_epochs), train_losses)
axs[0].plot(range(num_epochs), val_losses)
axs[0].legend(['Training Loss', 'Validation Loss'])
axs[0].set_title('Training and Validation DICE loss over epochs')

axs[1].plot(range(num_epochs), test_score)
axs[1].set_title('DICE score over epochs')

plt.show()
```



## 2.4 Training Procedure: Colorization Pre-training

Complete the rest of this problem in the cells below.

```
[ ]: ## Batch Size
train_batch_size = 10
validation_batch_size = 10

## Learning Rate/
learning_rate = 0.001

# Epochs (Consider setting high and implementing early stopping)
num_epochs = 50
```

```

[ ]: ## Augmented Image Transforms
augmented_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.RandomHorizontalFlip(),
    transforms.RandomVerticalFlip(),
])

## Image Dataloader
class ColorDataset(Dataset):

    """
    ImageDataset
    """

    def __init__(self,
                 input_dir,
                 op,
                 transforms=None):

        """
        ##TODO: Add support for colorization dataset

        Args:
            input_dir (str): Path to either colorization or segmentation
↪directory
            op (str): One of "train", "val", or "test" signifying the desired
↪split
            transforms (list or None): Image transformations to apply upon
↪loading.
        """
        self.transform = transforms
        self.op = op
        # Get directories
        try:
            if self.op == 'train':
                self.data_dir = os.path.join(input_dir, 'train_cor')
            elif self.op == 'val':
                self.data_dir = os.path.join(input_dir, 'validation_cor')
        except ValueError:
            print('op should be either train or val!')

    def __len__(self):
        """
        """
        return len(next(os.walk(self.data_dir))[1])-1

    def __getitem__(self,
                    idx):

```

```

        """

        """

        ## Load Image and Parse Properties
        idx = idx + 1
        img_name = str(idx) + '_gray.jpg'
        mask_name = str(idx) + '_input.jpg'
        img = io.imread(os.path.join(self.data_dir, str(idx), img_name))
        mask = io.imread(os.path.join(self.data_dir, str(idx), mask_name))
        if self.transform:
            img, mask = self.img_transform(img, mask)
        return img, mask

    def img_transform(self,
                      img,
                      mask):
        """

        """

        ## Apply Transformations to Image and Mask
        img = self.transform(img)
        mask = self.transform(mask)
        img = img.repeat(3,1,1)
        return img, mask

```

```

[ ]: ## Initialize your unet
n_classes = 3
model = UNET(n_classes)

## Initialize Dataloaders
train_dataset=ColorDataset(input_dir=colorization_data_dir, op="train",
    ↪transforms=augmented_transform)
validation_dataset=ColorDataset(input_dir=colorization_data_dir, op="val",
    ↪transforms=img_transform)
train_dataloader = DataLoader(train_dataset, batch_size=train_batch_size,
    ↪shuffle=True)
validation_dataloader = DataLoader(validation_dataset,
    ↪batch_size=validation_batch_size, shuffle=False)

## Initialize Optimizer and Learning Rate Scheduler
optimizer = torch.optim.Adam(model.parameters(),lr=learning_rate)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=10, gamma=0.1)
criterion = nn.MSELoss()
gpu_bool = torch.cuda.is_available()
if gpu_bool:
    model = model.cuda()

# Define lists to store training and validation losses

```

```

train_losses = []
val_losses = []

print("Start Training...")
for epoch in range(num_epochs):
    ##### Training #####
    print("\nEPOCH " +str(epoch+1)+" of "+str(num_epochs)+"\n")
    # TODO: Design your own training section
    model.train()
    train_loss = 0.0

    # Iterate through the training data
    for inputs, labels in train_dataloader:

        # Forward pass
        if gpu_bool:
            inputs = inputs.cuda()
            labels = labels.cuda()
        output = model.forward(inputs)

        # Compute loss
        loss = criterion.forward(output, labels)

        # Backpropagation
        optimizer.zero_grad() # Zero the gradients
        loss.backward()
        optimizer.step()

        # Update running training loss
        train_loss += loss.item()

    # Compute average training loss for the epoch
    avg_train_loss = train_loss / len(train_dataloader)
    train_losses.append(avg_train_loss)

    ##### Validation #####
    # TODO: Design your own validation section
    model.eval()
    val_loss = 0.0
    with torch.no_grad():
        for inputs, labels in validation_dataloader:
            if gpu_bool:
                inputs = inputs.cuda()
                labels = labels.cuda()
            output = model(inputs)
            loss = criterion.forward(output, labels)
            val_loss += loss.item()

```

```
avg_val_loss = val_loss / len(validation_dataloader)
val_losses.append(avg_val_loss)
scheduler.step()

print(f'Epoch [{epoch+1}/{num_epochs}], Train Loss: {avg_train_loss:.4f}, Val Loss: {avg_val_loss:.4f}')
```

Start Training...

EPOCH 1 of 50

Epoch [1/50], Train Loss: 0.0356, Val Loss: 0.0290

EPOCH 2 of 50

Epoch [2/50], Train Loss: 0.0353, Val Loss: 0.0286

EPOCH 3 of 50

Epoch [3/50], Train Loss: 0.0353, Val Loss: 0.0287

EPOCH 4 of 50

Epoch [4/50], Train Loss: 0.0353, Val Loss: 0.0290

EPOCH 5 of 50

Epoch [5/50], Train Loss: 0.0352, Val Loss: 0.0294

EPOCH 6 of 50

Epoch [6/50], Train Loss: 0.0352, Val Loss: 0.0288

EPOCH 7 of 50

Epoch [7/50], Train Loss: 0.0352, Val Loss: 0.0288

EPOCH 8 of 50

Epoch [8/50], Train Loss: 0.0352, Val Loss: 0.0288

EPOCH 9 of 50

Epoch [9/50], Train Loss: 0.0352, Val Loss: 0.0287

EPOCH 10 of 50

Epoch [10/50], Train Loss: 0.0352, Val Loss: 0.0288

EPOCH 11 of 50

Epoch [11/50], Train Loss: 0.0352, Val Loss: 0.0285

EPOCH 12 of 50

Epoch [12/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 13 of 50

Epoch [13/50], Train Loss: 0.0351, Val Loss: 0.0286

EPOCH 14 of 50

Epoch [14/50], Train Loss: 0.0351, Val Loss: 0.0286

EPOCH 15 of 50

Epoch [15/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 16 of 50

Epoch [16/50], Train Loss: 0.0351, Val Loss: 0.0286

EPOCH 17 of 50

Epoch [17/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 18 of 50

Epoch [18/50], Train Loss: 0.0351, Val Loss: 0.0287

EPOCH 19 of 50

Epoch [19/50], Train Loss: 0.0352, Val Loss: 0.0285

EPOCH 20 of 50

Epoch [20/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 21 of 50

Epoch [21/50], Train Loss: 0.0351, Val Loss: 0.0286

EPOCH 22 of 50



Epoch [22/50], Train Loss: 0.0351, Val Loss: 0.0286

EPOCH 23 of 50

Epoch [23/50], Train Loss: 0.0351, Val Loss: 0.0286

EPOCH 24 of 50

Epoch [24/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 25 of 50

Epoch [25/50], Train Loss: 0.0351, Val Loss: 0.0286

EPOCH 26 of 50

Epoch [26/50], Train Loss: 0.0351, Val Loss: 0.0286

EPOCH 27 of 50

Epoch [27/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 28 of 50

Epoch [28/50], Train Loss: 0.0352, Val Loss: 0.0287

EPOCH 29 of 50

Epoch [29/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 30 of 50

Epoch [30/50], Train Loss: 0.0351, Val Loss: 0.0286

EPOCH 31 of 50

Epoch [31/50], Train Loss: 0.0351, Val Loss: 0.0286

EPOCH 32 of 50

Epoch [32/50], Train Loss: 0.0352, Val Loss: 0.0287

EPOCH 33 of 50

Epoch [33/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 34 of 50

Epoch [34/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 35 of 50

Epoch [35/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 36 of 50

Epoch [36/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 37 of 50

Epoch [37/50], Train Loss: 0.0352, Val Loss: 0.0285

EPOCH 38 of 50

Epoch [38/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 39 of 50

Epoch [39/50], Train Loss: 0.0351, Val Loss: 0.0286

EPOCH 40 of 50

Epoch [40/50], Train Loss: 0.0351, Val Loss: 0.0285

EPOCH 41 of 50

Epoch [41/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 42 of 50

Epoch [42/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 43 of 50

Epoch [43/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 44 of 50

Epoch [44/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 45 of 50

Epoch [45/50], Train Loss: 0.0351, Val Loss: 0.0286

EPOCH 46 of 50

Epoch [46/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 47 of 50

Epoch [47/50], Train Loss: 0.0352, Val Loss: 0.0286

EPOCH 48 of 50

Epoch [48/50], Train Loss: 0.0351, Val Loss: 0.0286

EPOCH 49 of 50

Epoch [49/50], Train Loss: 0.0351, Val Loss: 0.0286


EPOCH 50 of 50

Epoch [50/50], Train Loss: 0.0351, Val Loss: 0.0286

```
[ ]: fig = plt.figure(figsize=(7, 5)) # create the canvas for plotting
ax = plt.subplot()
ax.plot(train_losses, label = "Training Loss")
ax.plot(val_losses, label = "Validation Loss")
handles, labels = ax.get_legend_handles_labels()

ax.set_ylabel("Loss")
ax.set_xlabel("Epochs")
ax.legend(handles, labels)

plt.show()
```






```
[ ]: torch.save(model.state_dict(), f'{data_dir}/color_pretrained.dict')
model = UNET(3)
model.load_state_dict(torch.load(f'{data_dir}/color_pretrained.dict',
    ↪map_location=torch.device('cpu')))
train_dataset=ColorDataset(input_dir=colorization_data_dir, op="train",
    ↪transforms=img_transform)
```

```
[ ]: model.eval()
plt.figure()
img, label = train_dataset[2]
plt.imshow(label.cpu().permute(1,2,0).numpy())
plt.figure()
img, label = train_dataset[2]
plt.imshow(img.cpu().permute(1,2,0).numpy())
with torch.no_grad():
    plt.figure()

    plt.imshow(img.cpu().permute(1,2,0).numpy()+model(torch.unsqueeze(img,0)).
    ↪squeeze().permute(1, 2, 0).cpu().numpy())
```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).





## 2.5 Use Color Pretrained Model on Segmentation dataset

```
[ ]: n_classes = 6
model.out_layer = nn.Conv2d(32, n_classes, kernel_size = 1, stride = 1, padding=
    ↪= 0, bias = True)
## Batch Size
train_batch_size = 10
validation_batch_size = 10

## Learning Rate/
learning_rate = 0.001

# Epochs (Consider setting high and implementing early stopping)
num_epochs = 100
```

```
[ ]: ## Initialize your unet
n_classes = 6
model = UNET(n_classes)
```

```

## Initialize Dataloaders
train_dataset=ImageDataset(input_dir=segmentation_data_dir, op="train",
    ↪mask_json_path=mask_json, transforms=augmented_transform)
validation_dataset=ImageDataset(input_dir=segmentation_data_dir, op="val",
    ↪mask_json_path=mask_json, transforms=img_transform)
test_dataset=ImageDataset(input_dir=segmentation_data_dir, op="test",
    ↪mask_json_path=mask_json, transforms=img_transform)
train_dataloader = DataLoader(train_dataset, batch_size=train_batch_size,
    ↪shuffle=True)
validation_dataloader = DataLoader(validation_dataset,
    ↪batch_size=validation_batch_size, shuffle=False)
test_dataloader = DataLoader(test_dataset, batch_size=1, shuffle=False)
## Initialize Optimizer and Learning Rate Scheduler
optimizer = torch.optim.Adam(model.parameters(),lr=learning_rate)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=10, gamma=0.1)
loss_func = DICELoss(n_classes)
model = model.cuda()

# Define lists to store training and validation losses
train_losses = []
val_losses = []
test_score = []

print("Start Training...")
for epoch in range(num_epochs):
    ##### Training #####
    print("\nEPOCH " +str(epoch+1)+" of "+str(num_epochs)+"\n")
    # TODO: Design your own training section
    model.train()
    train_loss = 0.0

    # Iterate through the training data
    for inputs, labels in train_dataloader:
        inputs = inputs.cuda()
        labels = labels.cuda()
        output = model.forward(inputs)

        loss = loss_func.forward(output, labels)
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    # Update running training loss
    train_loss += loss.item()

# Compute average training loss for the epoch
train_loss = train_loss / len(train_dataloader)

```

```

train_losses.append(train_loss)

##### Validation #####
# TODO: Design your own validation section
model.eval()
val_loss = 0.0
with torch.no_grad():
    for inputs, labels in validation_dataloader:
        inputs = inputs.cuda()
        labels = labels.cuda()
        output = model(inputs)
        loss = criterion.forward(output, labels)
        val_loss += loss.item()

val_loss = val_loss / len(validation_dataloader)
val_losses.append(val_loss)
test_score_i = dice_score_dataset(model, test_dataloader, n_classes,
↪use_gpu=gpu_bool)
test_score.append(test_score_i)
scheduler.step()

print(f'Epoch [{epoch+1}/{num_epochs}], Train Loss: {avg_train_loss:.4f},
↪Val Loss: {avg_val_loss:.4f}')
print(f'Test Score: {test_score_i:.4f}')

```

Start Training...

EPOCH 1 of 100

Epoch [1/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.3601

EPOCH 2 of 100

Epoch [2/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.3434

EPOCH 3 of 100

Epoch [3/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.3954

EPOCH 4 of 100

Epoch [4/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.5849



EPOCH 5 of 100

Epoch [5/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6144

EPOCH 6 of 100

Epoch [6/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.5524

EPOCH 7 of 100

Epoch [7/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.3340

EPOCH 8 of 100

Epoch [8/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.5642

EPOCH 9 of 100

Epoch [9/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.5727

EPOCH 10 of 100

Epoch [10/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6628

EPOCH 11 of 100

Epoch [11/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6200

EPOCH 12 of 100

Epoch [12/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6557

EPOCH 13 of 100

Epoch [13/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6778

EPOCH 14 of 100

Epoch [14/100], Train Loss: 0.0351, Val Loss: 0.0286

Test Score: 0.6730

EPOCH 15 of 100

Epoch [15/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6733

EPOCH 16 of 100

Epoch [16/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6765

EPOCH 17 of 100

Epoch [17/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6721

EPOCH 18 of 100

Epoch [18/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6697

EPOCH 19 of 100

Epoch [19/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6769

EPOCH 20 of 100

Epoch [20/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6777

EPOCH 21 of 100

Epoch [21/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6782

EPOCH 22 of 100

Epoch [22/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6765

EPOCH 23 of 100

Epoch [23/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6746

EPOCH 24 of 100

Epoch [24/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6771

EPOCH 25 of 100

Epoch [25/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6812

EPOCH 26 of 100

Epoch [26/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6760

EPOCH 27 of 100

Epoch [27/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6774

EPOCH 28 of 100

Epoch [28/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6812

EPOCH 29 of 100

Epoch [29/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6761

EPOCH 30 of 100

Epoch [30/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6809

EPOCH 31 of 100

Epoch [31/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6782

EPOCH 32 of 100

Epoch [32/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6733

EPOCH 33 of 100

Epoch [33/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6788

EPOCH 34 of 100

Epoch [34/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6773

EPOCH 35 of 100

Epoch [35/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6741

EPOCH 36 of 100

Epoch [36/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6798

EPOCH 37 of 100

Epoch [37/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6761

EPOCH 38 of 100

Epoch [38/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6798

EPOCH 39 of 100

Epoch [39/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6774

EPOCH 40 of 100

Epoch [40/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6749

EPOCH 41 of 100

Epoch [41/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6774

EPOCH 42 of 100

Epoch [42/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6776

EPOCH 43 of 100

Epoch [43/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6734

EPOCH 44 of 100

Epoch [44/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6811

EPOCH 45 of 100

Epoch [45/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6811

EPOCH 46 of 100

Epoch [46/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6778

EPOCH 47 of 100

Epoch [47/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6782

EPOCH 48 of 100

Epoch [48/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6792

EPOCH 49 of 100

Epoch [49/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6740

EPOCH 50 of 100

Epoch [50/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6734

EPOCH 51 of 100

Epoch [51/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6779

EPOCH 52 of 100

Epoch [52/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6815

EPOCH 53 of 100

Epoch [53/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6800

EPOCH 54 of 100

Epoch [54/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6765

EPOCH 55 of 100

Epoch [55/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6795

EPOCH 56 of 100

Epoch [56/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6814

EPOCH 57 of 100

Epoch [57/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6799

EPOCH 58 of 100

Epoch [58/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6794

EPOCH 59 of 100

Epoch [59/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6782

EPOCH 60 of 100

Epoch [60/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6794

EPOCH 61 of 100

Epoch [61/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6767

EPOCH 62 of 100

Epoch [62/100], Train Loss: 0.0351, Val Loss: 0.0286

Test Score: 0.6727

EPOCH 63 of 100

Epoch [63/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6782

EPOCH 64 of 100

Epoch [64/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6788

EPOCH 65 of 100

Epoch [65/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6777

EPOCH 66 of 100

Epoch [66/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6815

EPOCH 67 of 100

Epoch [67/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6771

EPOCH 68 of 100

Epoch [68/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6766

EPOCH 69 of 100

Epoch [69/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6798

EPOCH 70 of 100

Epoch [70/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6804

EPOCH 71 of 100

Epoch [71/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6800

EPOCH 72 of 100

Epoch [72/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6816

EPOCH 73 of 100

Epoch [73/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6780

EPOCH 74 of 100

Epoch [74/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6761

EPOCH 75 of 100

Epoch [75/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6783

EPOCH 76 of 100

Epoch [76/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6831

EPOCH 77 of 100

Epoch [77/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6766

EPOCH 78 of 100

Epoch [78/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6761

EPOCH 79 of 100

Epoch [79/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6769

EPOCH 80 of 100

Epoch [80/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6782

EPOCH 81 of 100

Epoch [81/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6814



EPOCH 82 of 100

Epoch [82/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6788

EPOCH 83 of 100

Epoch [83/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6813

EPOCH 84 of 100

Epoch [84/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6817

EPOCH 85 of 100

Epoch [85/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6771

EPOCH 86 of 100

Epoch [86/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6820

EPOCH 87 of 100

Epoch [87/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6809

EPOCH 88 of 100

Epoch [88/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6790

EPOCH 89 of 100

Epoch [89/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6811

EPOCH 90 of 100

Epoch [90/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6774

EPOCH 91 of 100

Epoch [91/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6783

EPOCH 92 of 100

Epoch [92/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6766

EPOCH 93 of 100

Epoch [93/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6807

EPOCH 94 of 100

Epoch [94/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6801

EPOCH 95 of 100

Epoch [95/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6799

EPOCH 96 of 100

Epoch [96/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6794

EPOCH 97 of 100

Epoch [97/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6804

EPOCH 98 of 100

Epoch [98/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6775

EPOCH 99 of 100

Epoch [99/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6762

EPOCH 100 of 100

Epoch [100/100], Train Loss: 0.0351, Val Loss: 0.0286  
Test Score: 0.6777

```
[ ]: max_score = max(test_score)
max_score_idx = test_score.index(max_score)
print(f'At Epoch {max_score_idx}, Test DICE score reaches its max as:␣
↳{max_score:.4f}')
```

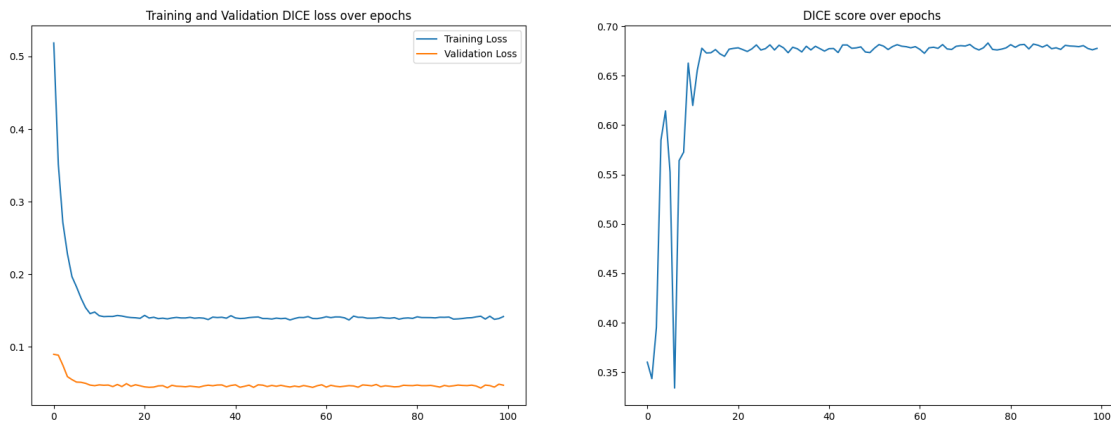
At Epoch 75, Test DICE score reaches its max as: 0.6831

```
[ ]: # Graph
fig, axs = plt.subplots(1, 2, figsize=(20, 7))

axs[0].plot(range(num_epochs), train_losses)
axs[0].plot(range(num_epochs), val_losses)
axs[0].legend(['Training Loss', 'Validation Loss'])
axs[0].set_title('Training and Validation DICE loss over epochs')

axs[1].plot(range(num_epochs), test_score)
axs[1].set_title('DICE score over epochs')


plt.show()
```



```
[ ]: fig = plt.figure() # create the canvas for plotting
ax = plt.subplot()
ax.plot(train_losses, label = "Training Loss")
ax.plot(val_losses, label = "Validation Loss")
handles, labels = ax.get_legend_handles_labels()

ax.set_ylabel("Loss")
ax.set_xlabel("Epochs")
ax.legend(handles, labels)

plt.show()
```



### 3 Problem 2: Transfer Learning

#### 3.0.1 Imports

```
[47]: ## Import VGG and FashionMNIST
from torchvision.models import vgg16
from torchvision.datasets import FashionMNIST
```

#### 3.0.2 Data Loading

```
[48]: ## Specify Batch Size
train_batch_size = 32
test_batch_size = 32

## Specify Image Transforms
img_transform = transforms.Compose([
    transforms.Resize((224,224)),
    transforms.Grayscale(3),
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
```

```

])

## Download Datasets
train_data = FashionMNIST('./data', transform=img_transform, download=True,
    ↪train=True)
test_data = FashionMNIST('./data', transform=img_transform, download=True,
    ↪train=False)

## Initialize Dataloaders
training_dataloader = DataLoader(train_data, batch_size=train_batch_size,
    ↪shuffle=True)
test_dataloader = DataLoader(test_data, batch_size=test_batch_size,
    ↪shuffle=False)

```

Downloading <http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-images-idx3-ubyte.gz>

Downloading <http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-images-idx3-ubyte.gz> to ./data/FashionMNIST/raw/train-images-idx3-ubyte.gz

100%| | 26421880/26421880 [00:01<00:00, 18023259.01it/s]

Extracting ./data/FashionMNIST/raw/train-images-idx3-ubyte.gz to  
./data/FashionMNIST/raw

Downloading <http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-labels-idx1-ubyte.gz>

Downloading <http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-labels-idx1-ubyte.gz> to ./data/FashionMNIST/raw/train-labels-idx1-ubyte.gz

100%| | 29515/29515 [00:00<00:00, 305220.98it/s]

Extracting ./data/FashionMNIST/raw/train-labels-idx1-ubyte.gz to  
./data/FashionMNIST/raw

Downloading <http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-images-idx3-ubyte.gz>

Downloading <http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-images-idx3-ubyte.gz> to  
./data/FashionMNIST/raw/t10k-images-idx3-ubyte.gz

100%| | 4422102/4422102 [00:00<00:00, 5588060.85it/s]

Extracting ./data/FashionMNIST/raw/t10k-images-idx3-ubyte.gz to  
./data/FashionMNIST/raw

Downloading <http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-labels-idx1-ubyte.gz>

Downloading <http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-labels-idx1-ubyte.gz> to  
./data/FashionMNIST/raw/t10k-labels-idx1-ubyte.gz

100% | 5148/5148 [00:00<00:00, 20466613.26it/s]

Extracting ./data/FashionMNIST/raw/t10k-labels-idx1-ubyte.gz to  
./data/FashionMNIST/raw

### 3.0.3 Model Initialization and Training/Fine-tuning

Complete the rest of the assignment in the notebook below.

```
[57]: model = vgg16(pretrained=False)
model.classifier[-1] = nn.Linear(in_features=4096, out_features=10)
model = model.cuda()

def initialize_weights(m):
    """
    Initialize the weights and biases of the model.
    """
    if isinstance(m, nn.Conv2d) or isinstance(m, nn.Linear):
        torch.nn.init.kaiming_uniform_(m.weight)
        torch.nn.init.uniform_(m.bias)

model.apply(initialize_weights)
```

```
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:223:
UserWarning: Arguments other than a weight enum or `None` for 'weights' are
deprecated since 0.13 and may be removed in the future. The current behavior is
equivalent to passing `weights=None`.
  warnings.warn(msg)
```

```
[57]: VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace=True)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace=True)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
  ceil_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU(inplace=True)
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): ReLU(inplace=True)
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
  ceil_mode=False)
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```

(13): ReLU(inplace=True)
(14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(15): ReLU(inplace=True)
(16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(18): ReLU(inplace=True)
(19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(20): ReLU(inplace=True)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(22): ReLU(inplace=True)
(23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(25): ReLU(inplace=True)
(26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(27): ReLU(inplace=True)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace=True)
(30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
)
(avgpool): AdaptiveAvgPool2d(output_size=(7, 7))
(classifier): Sequential(
  (0): Linear(in_features=25088, out_features=4096, bias=True)
  (1): ReLU(inplace=True)
  (2): Dropout(p=0.5, inplace=False)
  (3): Linear(in_features=4096, out_features=4096, bias=True)
  (4): ReLU(inplace=True)
  (5): Dropout(p=0.5, inplace=False)
  (6): Linear(in_features=4096, out_features=10, bias=True)
)
)

```

```

[ ]: model = model.cuda()

criterion = nn.CrossEntropyLoss()
lr = 0.001
optimizer = torch.optim.SGD(model.parameters(), lr = lr)
num_epochs = 10

# Define lists to store training and validation losses
train_losses = []
accuracy = []

print("Start Training...")
for epoch in range(num_epochs):

```

```

##### Training #####
print("\nEPOCH " +str(epoch+1)+" of "+str(num_epochs)+"\n")
# TODO: Design your own training section
model.train()
train_loss = 0.0

# Iterate through the training data
for inputs, labels in training_dataloader:
    inputs = inputs.cuda()
    labels = labels.cuda()
    output = model.forward(inputs)
    loss = criterion(output, labels)

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    train_loss += loss.item()

# Compute average training loss for the epoch
avg_train_loss = train_loss / len(training_dataloader)
train_losses.append(avg_train_loss)
print(f'Epoch [{epoch+1}/{num_epochs}], Train Loss: {avg_train_loss:.4f}')

##### Test #####
# TODO: Design your own Test section
model.eval()
test_loss = 0.0
test_correct = 0
with torch.no_grad():
    for inputs, labels in test_dataloader:
        if gpu_bool:
            inputs = inputs.cuda()
            labels = labels.cuda()
        output = model(inputs)
        loss = criterion(output, labels)
        test_loss += loss.item()
        predict = output.argmax(axis=1)
        test_correct += (predict == labels).float().sum()

avg_test_loss = test_loss / len(test_dataloader.dataset)
test_accuracy = test_correct / len(test_dataloader.dataset)
accuracy.append(test_accuracy.item())
print(f'Test Loss: {avg_test_loss:.4f}')
print(f'Test Accuracy: {test_accuracy:.4f}')

```



Start Training..

EPOCH 1 of 10

Epoch [1/10], Train Loss: 1.0280

Test Loss: 0.0139

Test Accuracy: 0.8349

EPOCH 2 of 10

Epoch [2/10], Train Loss: 0.5088

Test Loss: 0.0122

Test Accuracy: 0.8634

EPOCH 3 of 10

Epoch [3/10], Train Loss: 0.4306

Test Loss: 0.0110

Test Accuracy: 0.8715

EPOCH 4 of 10

```
[ ]: fig = plt.figure(figsize=(7, 5)) # create the canvas for plotting
ax = plt.subplot()
ax.plot(train_losses, label = "Training Loss")
ax.plot(accuracy, label = "Test Accuracy")
handles, labels = ax.get_legend_handles_labels()
ax.set_xlabel("Epochs")
ax.legend(handles, labels)
```

### 3.0.4 Using Pretrained VGG16

```
[56]: model = vgg16(pretrained = True)
model.classifier[-1] = nn.Linear(in_features=4096, out_features=10)
for param in model.parameters():
    param.requires_grad = False
for param in model.classifier[-1].parameters():
    param.requires_grad = True
initialize_weights(model.classifier[-1])
```

```
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:208:
UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be
removed in the future, please use 'weights' instead.
```

```
warnings.warn(
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:223:
UserWarning: Arguments other than a weight enum or `None` for 'weights' are
```

deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing `weights=VGG16_Weights.IMAGENET1K_V1`. You can also use `weights=VGG16_Weights.DEFAULT` to get the most up-to-date weights.

`warnings.warn(msg)`

Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to  
/root/.cache/torch/hub/checkpoints/vgg16-397923af.pth  
100%| | 528M/528M [00:10<00:00, 50.3MB/s]

```
[ ]: model = model.cuda()

criterion = nn.CrossEntropyLoss()
lr = 0.001
optimizer = torch.optim.SGD(model.parameters(), lr = lr)
num_epochs = 10

# Define lists to store training and validation losses
train_losses = []
accuracy = []

print("Start Training...")
for epoch in range(num_epochs):
    ##### Training #####
    print("\nEPOCH " + str(epoch+1) + " of " + str(num_epochs) + "\n")
    # TODO: Design your own training section
    model.train()
    train_loss = 0.0

    # Iterate through the training data
    for inputs, labels in training_dataloader:
        inputs = inputs.cuda()
        labels = labels.cuda()
        output = model.forward(inputs)
        loss = criterion(output, labels)

        # Backpropagation
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    train_loss += loss.item()

    # Compute average training loss for the epoch
    avg_train_loss = train_loss / len(training_dataloader)
    train_losses.append(avg_train_loss)
    print(f'Epoch [{epoch+1}/{num_epochs}], Train Loss: {avg_train_loss:.4f}')
```

```

##### Test #####
# TODO: Design your own Test section
model.eval()
test_loss = 0.0
test_correct = 0
with torch.no_grad():
    for inputs, labels in test_dataloader:
        inputs = inputs.cuda()
        labels = labels.cuda()
        output = model(inputs)
        loss = criterion(output, labels)
        test_loss += loss.item()
        predict = output.argmax(axis=1)
        test_correct += (predict == labels).float().sum()

avg_test_loss = test_loss / len(test_dataloader.dataset)
test_accuracy = test_correct / len(test_dataloader.dataset)
accuracy.append(test_accuracy.item())
print(f'Test Loss: {avg_test_loss:.4f}')
print(f'Test Accuracy: {test_accuracy:.4f}')

```

Start Training...

EPOCH 1 of 10

Epoch [1/10], Train Loss: 0.9955  
 Test Loss: 0.0190  
 Test Accuracy: 0.7926

EPOCH 2 of 10

Epoch [2/10], Train Loss: 0.6827  
 Test Loss: 0.0173  
 Test Accuracy: 0.8063

EPOCH 3 of 10

Epoch [3/10], Train Loss: 0.6283  
 Test Loss: 0.0165  
 Test Accuracy: 0.8137

EPOCH 4 of 10

Epoch [4/10], Train Loss: 0.5976  
 Test Loss: 0.0160  
 Test Accuracy: 0.8193

EPOCH 5 of 10

Epoch [5/10], Train Loss: 0.5803

Test Loss: 0.0156

Test Accuracy: 0.8246

EPOCH 6 of 10

Epoch [6/10], Train Loss: 0.5627

Test Loss: 0.0153

Test Accuracy: 0.8271

EPOCH 7 of 10

Epoch [7/10], Train Loss: 0.5528

Test Loss: 0.0150

Test Accuracy: 0.8298

EPOCH 8 of 10

Epoch [8/10], Train Loss: 0.5411

Test Loss: 0.0148

Test Accuracy: 0.8305

EPOCH 9 of 10

Epoch [9/10], Train Loss: 0.5362

Test Loss: 0.0146

Test Accuracy: 0.8362

EPOCH 10 of 10

Epoch [10/10], Train Loss: 0.5284

Test Loss: 0.0145

Test Accuracy: 0.8367

```
[55]: print(train_losses)
```

```
[1.0672889280001323, 0.5154962116916975]
```

```
[ ]: fig = plt.figure() # create the canvas for plotting
      ax = plt.subplot()
      ax.plot(train_losses, label = "Training Loss")
      ax.plot(accuracy, label = "Test Accuracy")
      handles, labels = ax.get_legend_handles_labels()
      ax.set_xlabel("Epochs")
      ax.legend(handles, labels)
```

```
[ ]: <matplotlib.legend.Legend at 0x7fad7b5d3a90>
```

